

MIPROv2: Multiprompt Instruction PRoposal Optimizer Version 2

Overview

MIPROv2 is a **prompt optimizer** for DSPy programs that jointly optimizes both **instructions** and **few-shot demonstrations** for each module in a Language Model (LM) program. It uses **Bayesian Optimization** to efficiently search over the space of possible prompt configurations.

The core insight is that optimizing LM programs requires solving two interrelated challenges:

1. **The Proposal Problem:** How to efficiently generate high-quality candidate instructions and demonstrations
 2. **The Credit Assignment Problem:** How to determine which prompt configurations contribute to overall performance when you only have task-level feedback
-

The Three-Step Pipeline

MIPROv2 operates in three main phases:

Step 1: Bootstrap Few-Shot Examples

The algorithm generates candidate few-shot demonstrations through **rejection sampling**:

1. Randomly sample inputs $\langle x \rangle$ from the training set
2. Run them through the LM program $\langle \Phi(x) \rangle$ to generate traces
3. If the output scores well according to metric $\langle \mu \rangle$ (i.e., $\langle \mu(\Phi(x), x') \geq \lambda \rangle$), keep the input/output traces as valid demonstration candidates
4. Repeat until you have $\langle \text{num_candidates} \rangle$ sets of demonstrations

This creates multiple candidate demo sets, each containing:

- $\langle \text{max_bootstrapped_demos} \rangle$ bootstrapped examples (from successful traces)
- $\langle \text{max_labeled_demos} \rangle$ basic examples sampled directly from training

Step 2: Propose Instruction Candidates

For each module in the program, MIPROv2 generates diverse instruction candidates using a **grounded proposer**. The proposer LM receives:

Context Component	Description
Dataset Summary	Auto-generated description of patterns in the training data
Program Summary	Description of the program's control flow and purpose
Bootstrapped Demos	Example input/outputs showing successful task completions
Tips	Randomly sampled prompt engineering suggestions (e.g., "be creative", "be concise")

The tips encourage diversity in proposed instructions. Available tips include:

- `"none"`: No additional guidance
- `"creative"`: "Don't be afraid to be creative!"
- `"simple"`: "Keep the instruction clear and concise."
- `"description"`: "Make sure your instruction is very informative and descriptive."
- `"high_stakes"`: Include a high stakes scenario
- `"persona"`: Provide the LM with a relevant persona

Step 3: Bayesian Optimization Search

This is the key innovation. Instead of random search, MIPROv2 uses a **Tree-structured Parzen Estimator (TPE)** to learn which combinations of instructions and demonstrations work best:

```
for trial in 1..num_trials:
    # Sample a configuration using TPE's learned priors
    config = TPE.sample(
        instruction_for_module_1: choice from candidates,
        instruction_for_module_2: choice from candidates,
        demos_for_module_1: choice from candidate sets,
        demos_for_module_2: choice from candidate sets,
        ...
    )

    # Evaluate on a minibatch (not full dataset)
    score = evaluate(program_with_config, minibatch)

    # Update TPE's belief about good configurations
    TPE.update(config, score)

    # Periodically do full evaluation on best candidates
    if trial % minibatch_full_eval_steps == 0:
        full_evaluate(best_averaging_config)
```

Key advantages of this approach:

1. **Efficient exploration:** TPE learns to focus on promising regions of the search space
 2. **Credit assignment:** The Bayesian model learns which parameters (instructions, demos) matter most
 3. **Minibatching:** Evaluating on small batches allows more trials within the same budget
 4. **Joint optimization:** Instructions and demos are optimized together, capturing interactions
-

Algorithm Pseudocode

Algorithm: MIPROv2

Input:

- Program Φ with m modules
- Metric μ
- Training data D
- num_candidates (N)
- num_trials (T)
- minibatch_size (B)

Phase 1 - Bootstrap Demonstrations:

```
for each module i in  $\Phi$ :  
    demo_candidates[i] = []  
    for j in 1..N:  
        demo_set = bootstrap_demos( $\Phi$ ,  $D$ ,  $\mu$ , max_bootstrapped, max_labeled)  
        demo_candidates[i].append(demo_set)
```

Phase 2 - Propose Instructions:

```
for each module i in  $\Phi$ :  
    instruction_candidates[i] = []  
    dataset_summary = summarize_dataset( $D$ )  
    program_summary = summarize_program( $\Phi$ )  
    for j in 1..N:  
        tip = random_choice(tips)  
        demos = random_choice(demo_candidates[i])  
        instruction = proposer_LM(dataset_summary, program_summary, demos, tip)  
        instruction_candidates[i].append(instruction)
```

Phase 3 - Bayesian Optimization:

```
Initialize TPE with uniform priors over all candidates  
best_score =  $-\infty$   
best_config = None  
  
for trial in 1..T:  
    # Sample configuration using TPE
```

```

config = {}
for each module i:
    config[f"instruction_{i}"] = TPE.sample(instruction_candidates[i])
    config[f"demos_{i}"] = TPE.sample(demo_candidates[i])

# Evaluate on minibatch
minibatch = random_sample(D, B)
program_configured = apply_config(Φ, config)
score = evaluate(program_configured, minibatch, μ)

# Update TPE
TPE.update(config, score)

# Periodic full evaluation
if trial % minibatch_full_eval_steps == 0:
    best_averaging = get_best_averaging_config(TPE)
    full_score = evaluate(apply_config(Φ, best_averaging), D, μ)
    if full_score > best_score:
        best_score = full_score
        best_config = best_averaging

return apply_config(Φ, best_config)

```

Key Parameters

Parameter	Description	Default
metric	Evaluation function for program outputs	Required
auto	Preset mode: "light", "medium", "heavy"	"light"
num_candidates	Number of instruction/demo candidates per module	Set by auto
max_bootstrapped_demos	Max demos from successful traces	4
max_labeled_demos	Max demos directly from training	4
prompt_model	LM for proposing instructions	Default LM
task_model	LM used in the actual program	Default LM
minibatch_size	Size of evaluation batches	35
minibatch_full_eval_steps	Trials between full evaluations	5

Auto Modes

The `auto` parameter provides convenient presets:

Mode	Use Case	Trials	Candidates
"light"	Quick optimization, limited budget	Fewer	Fewer
"medium"	Balanced optimization	Medium	Medium
"heavy"	Maximum quality, larger budget	Many	Many

Proposer Configurations

MIPROv2 supports several flags to control what context the proposer sees:

Flag	Description
<code>program_aware_proposer</code>	Include program structure summary
<code>data_aware_proposer</code>	Include dataset characteristics summary
<code>tip_aware_proposer</code>	Include random prompt engineering tips
<code>fewshot_aware_proposer</code>	Include bootstrapped demonstrations

Comparison with Related Methods

Method	Instructions	Demos	Search Strategy	Credit Assignment
Bootstrap Random Search	✗	✓	Random	None
OPRO	✓	✗	History-based LM	LM implicit
Module-Level OPRO	✓	✗	Per-module history	Program score proxy
0-Shot MIPRO	✓	✗	Bayesian (TPE)	Surrogate model
Bayesian Bootstrap	✗	✓	Bayesian (TPE)	Surrogate model
MIPRO / MIPROv2	✓	✓	Bayesian (TPE)	Surrogate model

Key Findings from the Paper

The MIPRO paper (Opsahl-Ong et al., 2024) established several important lessons:

1. **Bootstrapped demonstrations are crucial:** For most tasks, optimizing demos alone outperforms optimizing instructions alone
 2. **Joint optimization is best:** MIPRO (instructions + demos together) generally achieves the highest performance
 3. **Instructions matter for conditional rules:** When tasks have subtle rules that can't be inferred from few examples, instruction optimization becomes essential
 4. **Grounding helps but varies:** The usefulness of dataset/program summaries varies by task—MIPRO++ can learn optimal proposal strategies
 5. **Bayesian optimization is robust:** The surrogate model handles noisy evaluations and efficiently explores the search space
-

Example Usage

```
python
```

```

import dspy
from dspy.teleprompt import MIPROv2

# Configure your LM
lm = dspy.LM('openai/gpt-4o-mini', api_key='...')
dspy.configure(lm=lm)

# Define your metric
def my_metric(prediction, example):
    return prediction.answer == example.answer

# Initialize optimizer
optimizer = MIPROv2(
    metric=my_metric,
    auto="medium", # or "light" / "heavy"
)
# Define your program
class MyProgram(dspy.Module):
    def __init__(self):
        self.generate = dspy.ChainOfThought("question -> answer")

    def forward(self, question):
        return self.generate(question=question)

# Optimize
optimized_program = optimizer.compile(
    MyProgram(),
    trainset=train_data,
    # Optional: provide a validation set
    valset=val_data,
)
# Save for later use
optimized_program.save("optimized_program.json")

```

Practical Recommendations

1. Start with `auto="light"` to get a quick baseline, then increase to `"medium"` or `"heavy"` if you have compute budget
2. For zero-shot optimization (no few-shot demos), set `max_bootstrapped_demos=0` and `max_labeled_demos=0`

3. **For tasks with complex rules**, provide a good seed instruction that describes the rules, as optimizers struggle to infer all rules automatically
 4. **Use a stronger teacher model** for bootstrapping on challenging tasks (e.g., GPT-4 instead of the task model)
 5. **Monitor optimization progress**: The Bayesian model's learned importance scores can reveal which components matter most for your task
-

References

- Opsahl-Ong, K., Ryan, M. J., Purtell, J., et al. (2024). *Optimizing Instructions and Demonstrations for Multi-Stage Language Model Programs*. arXiv:2406.11695
- DSPy Documentation: <https://dspy.ai/api/optimizers/MIPROv2/>